Detecting AI-Generated Music Using LambdaResNet and Swin Transformers on Mel-Spectrograms

ST311 Group 12
Candidate Numbers: 44535, 40538

Motivation

Why Detect AI-Generated Music?

- AI music generators (e.g. Suno, MusicGen) can now mimic human composition with striking realism
- Threats to authenticity and copyright: Artists face risk of being mimicked or replaced
- Platforms may be flooded with synthetic content, risking monetization abuse and user trust

Music labels sue AI song generators Suno and Udio for copyright infringement

Software steals songs to 'spit out' similar tunes, lawsuit says, asking for $150,000 a work in compensation

Sony, Universal and Warner are suing AI song generators, alleging they are exploiting the copyrighted music of artists from Mariah Carey to Chuck Berry. Photograph: Damian Dovarganes/AP

AS SUNO AND UDIO ADMIT TRAINING AI WITH UNLICENSED MUSIC, RECORD INDUSTRY SAYS: 'THERE'S NOTHING FAIR ABOUT STEALING AN ARTIST'S LIFE'S WORK.'

AUGUST 5, 2024 BY DANIEL TENCER

Credit: Shutterstock

Research Question

Can 5-second mel-spectrograms be classified using small vision models?

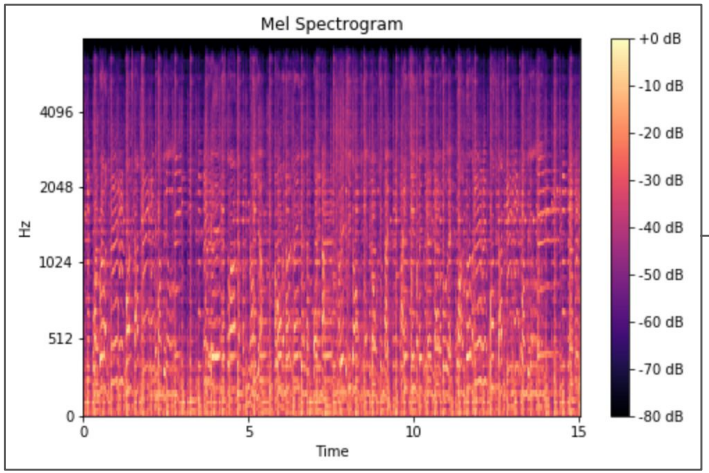

Input

Is the song AI-generated? (Yes/No)

Output

Dataset: SONICS

- Contains 50,000 real + 50,000 AI generated songs
- In our use case, we selected 10,000 real + 10,000 AI generated songs randomly, creating our own dataset
- Train / Val / Test ratio = 60:20:20

Model Architectures

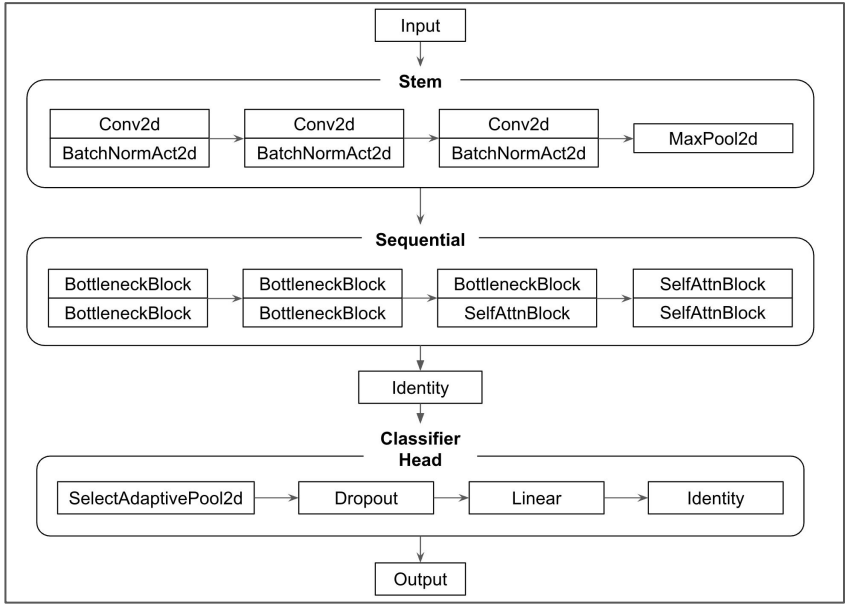

LambdaResNet26rp_256 (~11M params)

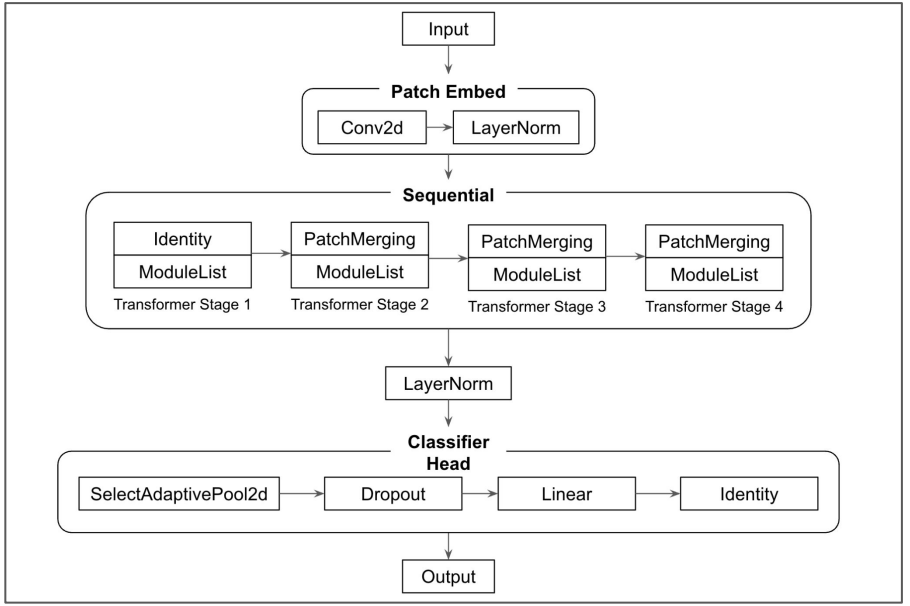

Swin Transformer V2 (~51M params)

Algorithm Pipeline

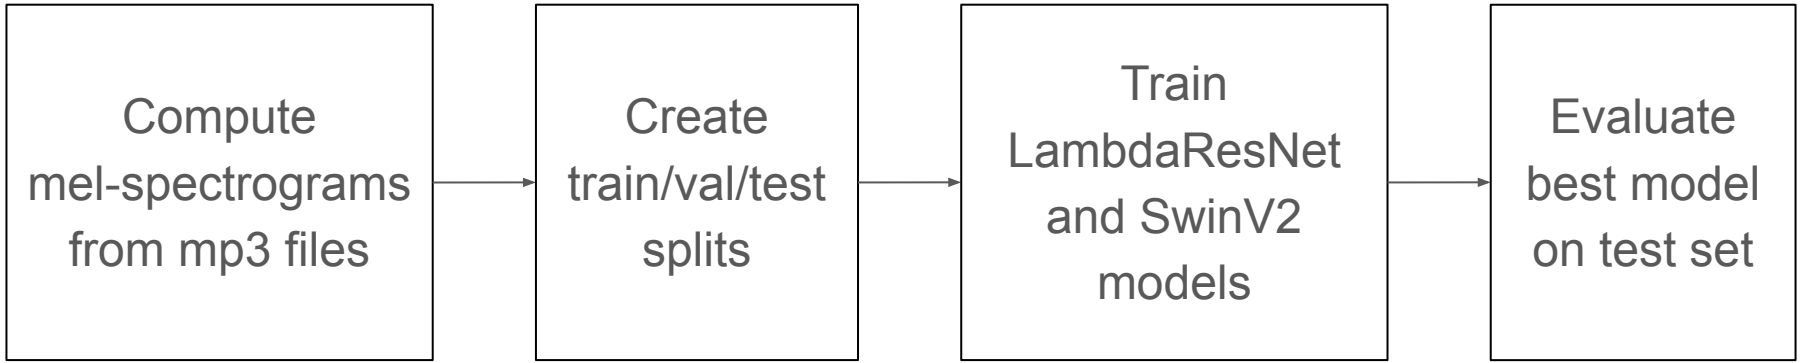

Results

Model	Precision	Recall	F1 Score	Specificity	AUC-ROC
LambdaResNet26rp_256	0.9910	0.9925	0.9918	0.9910	0.9996
Swin Transformer V2 Small	**0.9995**	**0.9990**	**0.9992**	**0.9995**	**1.0000**

- **Swin Transformer** achieves **near-perfect classification** across all metrics:
  - 99.9%+ precision → Very few false positives
  - AUC-ROC = 1.000 → Perfect separation between classes
- **LambdaResNet** is slightly **less accurate** but still exceptional:
  - Over 99% across all metrics
  - Best for speed and efficiency (fewer params, faster inference)
- Both models are well-calibrated with balanced specificity and recall
- Swin Transformer is ideal for high-stakes, high-accuracy settings
- LambdaResNet is better for real-time or edge deployment

Limitations

- Used only SONICS dataset (might be biases in music genre)
- Binary classification only
- Trained on shorter music clips

Future Work

- Multi-class detection
- Longer audio modelling (using full songs)
- More explainability (why the model classified a clip as real / fake)